

Conversational AI: Speech Recognition & Synthesis

Assignment

- This work, "Speech Commands: Dataset for Limited-Vocabulary Speech Recognition", presents a dataset of 65,000 one-second utterances for limited voice command recognition tasks. The dataset has been used more as a standard testbed to validate voice recognition models and to make the creation and evaluation of small-footprint models in resource-constrained environments easier. The dataset includes 30 short words, silence, and background noise.

Summary of my work:-

- **Data Preparation:** Labels and Features: For picture labels and features, use numpy arrays.
- **One-Hot Encoding:** Use `tf.keras.utils.to_categorical` to convert categorical labels into a format that is one-hot encoded.
- **Creation of Datasets:**
- **TensorFlow Dataset:** Use `tf.data` to transform numpy arrays into TensorFlow datasets. `From_tensor_slices` dataset.

To improve training dynamics, apply batching and shuffling to the training dataset.

- **Model Defined:**
- **Architecture:** Construct a Sequential model using `MaxPooling2D`, `Dense`, `Flatten`, and `Conv2D` layers.
- **Output Layer:** Verify that there are the appropriate number of units (e.g., 12 units for 12 classes) in the output layer.

- **Gathering:**
- **Diminished Function:** For labels that are encoded one-hot, use `categorical_crossentropy`.
- **Optimizer:** Use a suitable optimizer, such as Adam, to compile the model and provide assessment metrics.
- **Instruction:**

Model Fitting: Use a model to train the model in line with the instruction and verification along with test data